

Neural Networks and Learning Systems  
TBMI26 / 732A55  
2019

**Lecture 3**

**Supervised learning – Neural networks**

Magnus Borga  
magnus.borga@liu.se

### Recap - Supervised learning


- **Task:** Learn to predict/classify new data from labeled examples.
- **Input:** Training data examples  $\{\mathbf{x}_i, y_i\} i=1\dots N$ , where  $\mathbf{x}_i$  is a feature vector and  $y_i$  is a class label.
- **Output:** A function  $f(\mathbf{x}; w_1, \dots, w_k)$  that can predict the class label of  $\mathbf{x}$ .

Find a function  $f$  and adjust the parameters  $w_1, \dots, w_k$  so that new feature vectors are classified correctly. Generalization!!


2

### Linear separability

Linearly separable




Non-linearly separable



3

### The XOR problem

$x_1$	$x_2$	$f$
0	0	0
0	1	1
1	0	1
1	1	0




Not linearly separable!

4

## Multi-layer perceptron

a.k.a Neural Network



5

## History of neural networks




- 1960's: Large enthusiasm around the perceptron and "connectionism" (Frank Rosenblatt).
- 1969: Limitations of the perceptron made clear in a paper by Minsky & Papert, e.g., the XOR problem.
- "Winter period" – little research
- 1980's: Revival of connectionism and neural networks:
  - Multi-layer networks can solve nonlinear problems (this was known before, but not how to train them!)
  - Back-propagation training algorithm
- 1990's: Reduced interest, other methods seemed more promising
- 2010's: Renewed interest – "Deep learning"

6


## A simple 1D example

Training samples with only one feature value!




7

## Example solution



## Nonlinear mapping to a new feature space!



## Key: The hidden layer(s)

- The output layer requires linear separability. The purpose of the hidden layers is to make the problem linearly separable!
- Cover's theorem (1965):** The probability that classes are linearly separable increases when the features are nonlinearly mapped to a higher-dimensional feature space.

10


## The Perceptron Revisited




Minimize the following cost function

$$\mathcal{E}(\mathbf{w}) = \sum_{i=1}^N (\sigma(\mathbf{w}^T \mathbf{x}_i) - y_i)^2$$

$$\sigma(\mathbf{w}^T \mathbf{x}) = \mathbf{w}^T \mathbf{x}$$



$$\sigma(\mathbf{w}^T \mathbf{x}) = \tanh(\mathbf{w}^T \mathbf{x})$$




11

## Nonlinear activation functions


- Step/sign function

Not differentiable – cannot be optimized!  
(by gradient search)




- Hyperbolic tangent

$$\sigma(s) = \tanh(s) \quad \sigma' = 1 - \tanh^2(s) = 1 - \sigma^2$$




- The Fermi-function

$$\sigma(s) = \frac{1}{1 + e^{-s}} \quad \sigma' = \sigma(1 - \sigma)$$



- Gaussian function


$$\sigma(s; \gamma) = e^{-\frac{s^2}{\gamma^2}} \quad \sigma'(s; \gamma) = -\frac{2s}{\gamma} \sigma$$



12

## Example - Radial Basis Functions

For example a Gaussian



13

## Updated minimization algorithm

$$\mathcal{E}(\mathbf{w}) = \sum_{i=1}^N (\sigma(\mathbf{w}^T \mathbf{x}_i) - y_i)^2$$

$$\frac{\partial \mathcal{E}}{\partial \mathbf{w}} = 2 \sum_{i=1}^N (\sigma(\mathbf{w}^T \mathbf{x}_i) - y_i) \sigma'(\mathbf{w}^T \mathbf{x}_i) \mathbf{x}_i$$

Gradient descent:

$$\mathbf{w}_{t+1} = \mathbf{w}_t - \eta \frac{\partial \mathcal{E}}{\partial \mathbf{w}} \quad (\text{Eq. 1})$$


**Algorithm:**

1. Start with a random  $\mathbf{w}$
2. Iterate Eq. 1 until convergence

14


## Example

$$\sigma(s) = s$$




Same as in lecture 2!

$$\sigma(s) = \tanh(s)$$



15

## Training multi-layer neural networks



**Cost function**

$$\mathcal{E}(\mathbf{w}) = \sum_{k=1}^K \sum_{m=1}^M [y_{mk} - z_{mk}(\mathbf{w})]^2$$

Annotations:   
 -  $\mathcal{E}(\mathbf{w})$ : all weights   
 -  $K$ : # training examples   
 -  $M$ : # output nodes   
 -  $y_{mk}$ : desired output   
 -  $z_{mk}(\mathbf{w})$ : actual output

16


## Stochastic gradient descent

Update using one (K=1) training example

$$\mathcal{E}(\mathbf{w}) = \sum_{m=1}^M [y_m - z_m(\mathbf{w})]^2$$

$$w_{ij}^{t+1} = w_{ij}^t - \eta \frac{\partial \mathcal{E}}{\partial w_{ij}}$$

From node i to node j  
in a layer



17

## The chain rule

$$f(g(x))$$

$$\frac{\partial f}{\partial x} = \frac{\partial f}{\partial g} \frac{\partial g}{\partial x}$$

$$f(g(x), h(x))$$

$$\frac{\partial f}{\partial x} = \frac{\partial f}{\partial g} \frac{\partial g}{\partial x} + \frac{\partial f}{\partial h} \frac{\partial h}{\partial x}$$

Examples:

$$f(x) = \sin(x^2) \quad \frac{\partial f}{\partial x} = \cos(x^2) 2x$$

$$f(x) = x^2 \sin(x) \quad \frac{\partial f}{\partial x} = 2x \sin(x) + x^2 \cos(x)$$


18

## The error back propagation algorithm

- an exercise of the chain rule!

$$\frac{\partial \mathcal{E}}{\partial w_{ij}} = \frac{\partial \mathcal{E}}{\partial z_j} \frac{\partial z_j}{\partial w_{ij}} = \frac{\partial \mathcal{E}}{\partial z_j} \frac{\partial z_j}{\partial s_j} \frac{\partial s_j}{\partial w_{ij}}$$

$$\mathcal{E}(\mathbf{w}) = \sum_{m=1}^M [y_m - z_m(\mathbf{w})]^2$$



19

## Back propagation, cont.

$$\mathcal{E}(\mathbf{w}) = \sum_{m=1}^M [y_m - z_m(\mathbf{w})]^2 \quad \frac{\partial \mathcal{E}}{\partial w_{ij}} = \frac{\partial \mathcal{E}}{\partial z_j} \frac{\partial z_j}{\partial s_j} \frac{\partial s_j}{\partial w_{ij}}$$

$$\frac{\partial \mathcal{E}}{\partial z_j} = -2(y_j - z_j)$$


$$\frac{\partial z_j}{\partial s_j} = \sigma'(s_j) = 1 - \sigma(s_j)^2 = 1 - z_j^2 \quad \text{If } \sigma(s) = \tanh(s) \text{ is used!}$$

$$\frac{\partial s_j}{\partial w_{ij}} = h_i \quad (\text{input } i \text{ to unit } j)$$

20

## Updating the hidden layer

$$\frac{\partial \varepsilon}{\partial v_{ij}} = ?$$



$$\varepsilon(\mathbf{v}) = \sum_{m=1}^M [y_m - z_m(\mathbf{v})]^2$$

A weight in the hidden layer affects **all** output nodes!

$$\varepsilon(z_1(\mathbf{v}), \dots, z_M(\mathbf{v}))$$

$$\frac{\partial \varepsilon}{\partial v_{ij}} = \sum_{k=1}^M \frac{\partial \varepsilon}{\partial z_k} \frac{\partial z_k}{\partial v_{ij}} = \dots \quad \text{Exercise!}$$

Chain rule!

Continue expanding!

21

## Back propagation – Summary


- Two phases:
  - Forward propagation: Propagate a training example through the network
  - Backward propagation: Propagate the error relative the desired output backwards in the net and update parameter weights.
- Batch update: update after all examples have been presented.

$$\Delta w_{ij} = -\eta \sum_{k=1}^K \frac{\partial \varepsilon(k)}{\partial w_{ij}}$$

22

## Decision boundaries

Neural networks can produce very complex class boundaries!




$$f(x_1, x_2) = \text{sign}(z(x_1, x_2))$$



23

## Reminder - Magic is not possible!

No neural network, however complex, can separate inseparable classes!



Finding and extracting suitable features are the critical problems in machine learning!

24


## Pros and cons of neural networks

- A multi-layer neural network can learn any class boundaries.
- The large number of parameters is a problem:
  - Local optima  $\rightarrow$  suboptimal performance
  - Overfitting  $\rightarrow$  poor generalization
  - Slow convergence  $\rightarrow$  long training times

25

## Overfitting


- The large number of parameters makes it possible to produce overly complicated boundaries.



- A too good fit to the training data can perform poorly for new cases, i.e. worse generalization properties!

26


## Example



27

## Over-fitting


- The error on training data *always* decreases with increased training
- The error on validation data (the generalization error) decreases in the beginning, but can then start to increase if over-fitting occurs!



28

## Preventing overfitting in neural networks

- **Early stopping:**  
Pause training regularly and calculate the performance on the validation data.
- Caution: Validation data becomes training data! Will bias evaluation.
- That's why we need a third dataset for testing – the test data



29

## Faster convergence

- Normalize input features, e.g. to the range  $[-1,1]$ .
- Separate and adaptive step length  $\eta$  for each weight:
  - If the derivative has the same sign in several consecutive steps,  $\eta$  should increase. If the derivative change sign,  $\eta$  should decrease.
- Introduce a *momentum term*:

$$\Delta w_{ji}(t) = \underbrace{\alpha \Delta w_{ji}(t-1)}_{\text{A part of the previous update}} - \eta \frac{\partial \varepsilon(t)}{\partial w_{ji}}$$

30

## How many layers?

- 1 hidden layer is enough to produce any classification boundary.
- Complex boundaries more compactly obtained with many non-linear layers - less nodes in total compared to 1-layer solution.
- With ordinary 'backprop' training, no performance advantage with many hidden layers.
  - Vanishing gradient problem:  
Error gradients become very small for early layers in the network  $\rightarrow$  weights are not updated.
- Modifications of activation function etc. enable training in many layers  $\rightarrow$  "deep learning"!

31