

Neural Networks and Learning Systems
TBMI26 / 732A55
2019

Lecture 8

Unsupervised Learning – Dimensionality Reduction, Clustering

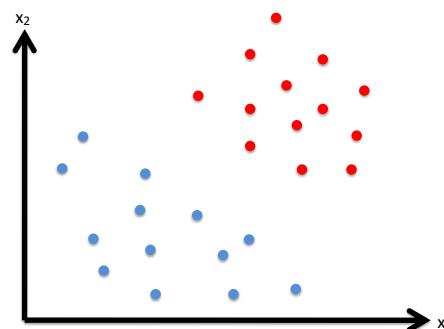
Magnus Borga
magnus.borga@liu.se

Three main categories of machine learning methods

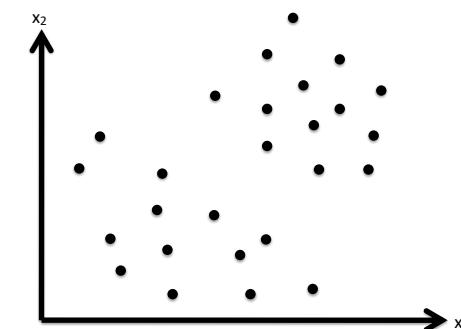
- **Supervised learning (predictive)**
Learn to generalize and classify new data based on labeled training data.
 - Pattern recognition
 - Classification
- **Reinforcement learning (active)**
Generate policies/strategies that lead to a (possibly delayed) reward. Learning by doing.
- **Unsupervised learning (descriptive)**
Discover structure and relationships in complex high-dimensional data.

2

Supervised learning – labeled samples

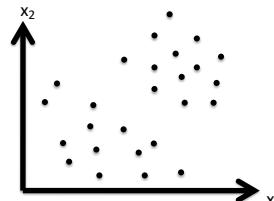


Unsupervised learning – unlabeled samples



Unsupervised learning

- **Task:** Find underlying structure in data.
- **Input:** Training data examples $\{x_i\}$ $i=1\dots N$.
- **Output:** Description of the data in a simpler form, e.g., with fewer dimensions or parameters.



5

Unsupervised learning

- Optimizes an internal cost function, e.g.
 - max variance (PCA)
 - max class separability (LDA)
- Finds a new representation of the data

6

Applications

- Feature extraction
 - find order or structure in data
- Dimensionality reduction
 - keep the most "important" parts of the signal

7

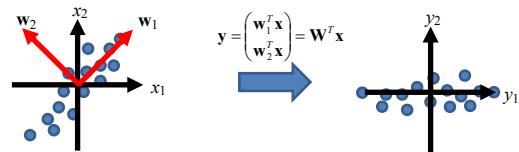
Too many dimensions/features

Correlated features or features that do not carry any information:

- Introduce noise in the analysis/classification
- Introduce more parameters in the learning model
 - More local optima in the optimization
 - Poorer generalization
 - Higher computational effort
- Difficult to visualize high-dimensional data

8

Projection / linear transformation



Can "uncorrelate" data through a linear transformation!!

$$\mathbf{C}_x = \begin{bmatrix} \text{Var}(x_1) & \text{Cov}(x_1, x_2) \\ \text{Cov}(x_2, x_1) & \text{Var}(x_2) \end{bmatrix}$$

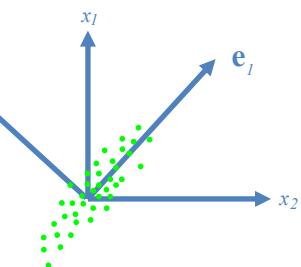
$$\mathbf{C}_y = \begin{bmatrix} \text{Var}(y_1) & 0 \\ 0 & \text{Var}(y_2) \end{bmatrix}$$

9

PCA

Principal Component Analysis

- Pearson 1901
(A.k.a. Hotelling-transform or Karhunen-Loéve-transform)
- Coordinate transformation to an orthogonal basis where the data is uncorrelated.
- Dimensionality reduction that preserves maximum variance (minimizes the mean square error).



10

Maximize the variance

The variance in direction $\hat{\mathbf{w}}$: (Suppose \mathbf{x} has mean 0.)

Normalized vector

$$\sigma_{\hat{\mathbf{w}}}^2 = E[(\mathbf{x}^T \hat{\mathbf{w}})^2] = E[(\hat{\mathbf{w}}^T \mathbf{x})(\mathbf{x}^T \hat{\mathbf{w}})]$$

$$= \hat{\mathbf{w}}^T E[\mathbf{x} \mathbf{x}^T] \hat{\mathbf{w}} = \hat{\mathbf{w}}^T \mathbf{C} \hat{\mathbf{w}} = \frac{\mathbf{w}^T \mathbf{C} \mathbf{w}}{\mathbf{w}^T \mathbf{w}}$$

The covariance matrix of \mathbf{x} .

11

Maximize the variance

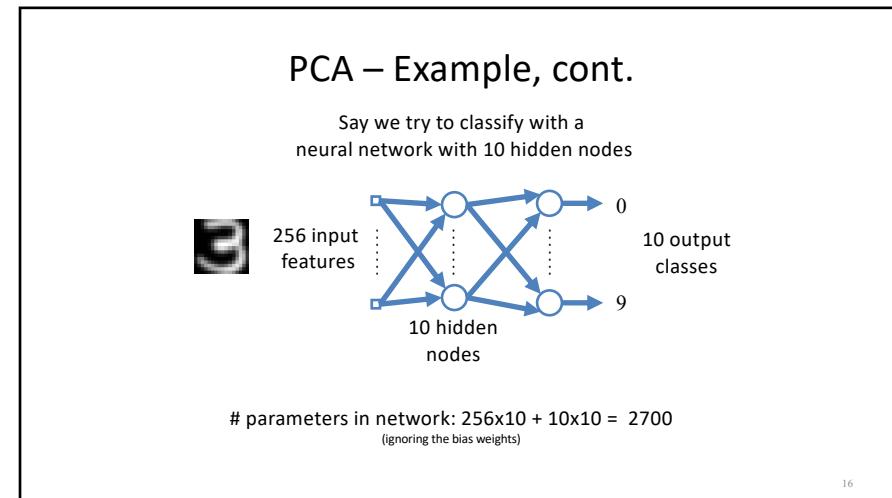
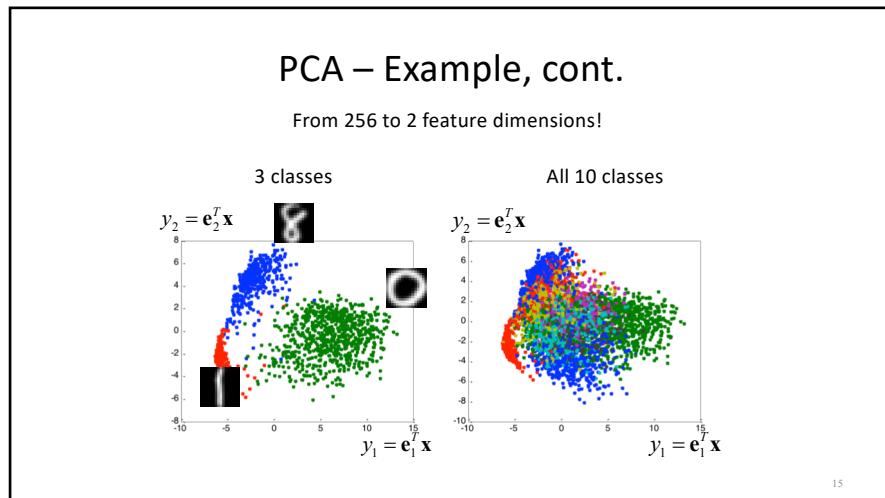
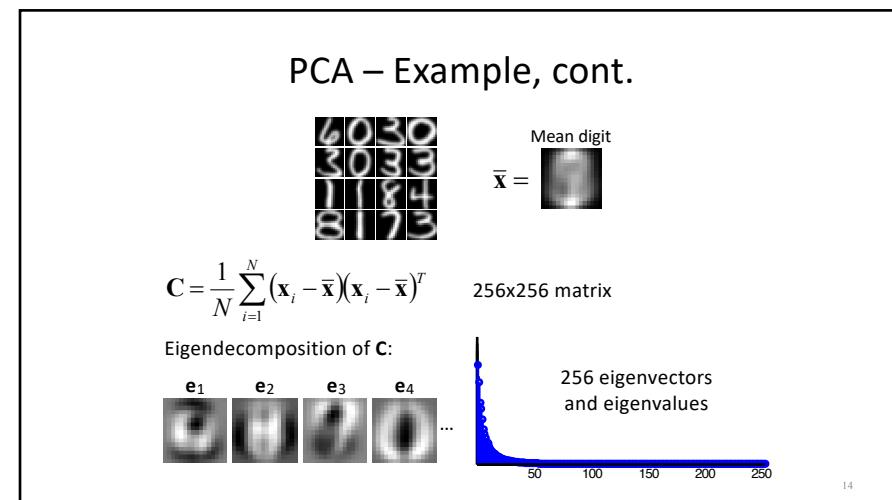
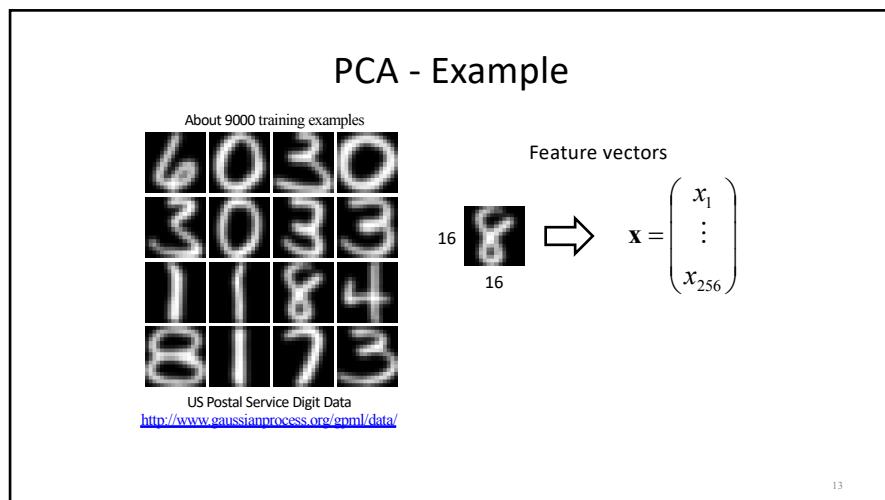
$$\sigma_{\hat{\mathbf{w}}}^2 = \frac{\mathbf{w}^T \mathbf{C} \mathbf{w}}{\mathbf{w}^T \mathbf{w}}$$

$$\frac{\partial \sigma_{\hat{\mathbf{w}}}^2}{\partial \mathbf{w}} = \frac{2}{\mathbf{w}^T \mathbf{w}} (\mathbf{C} \mathbf{w} - \sigma_{\hat{\mathbf{w}}}^2 \mathbf{w}) = 0 \Rightarrow$$

$$\mathbf{C} \mathbf{w} = \sigma_{\hat{\mathbf{w}}}^2 \mathbf{w}$$

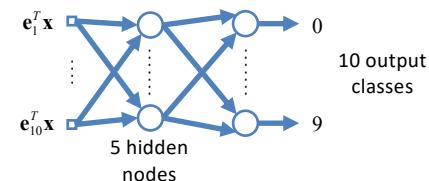
PCA is the Eigen-value decomposition of the data covariance matrix.

12



PCA – Example, cont.

If we reduce the input dimensionality first, we may be able to do the classification with a smaller network, e.g., 10 principal components as input and 5 hidden nodes.

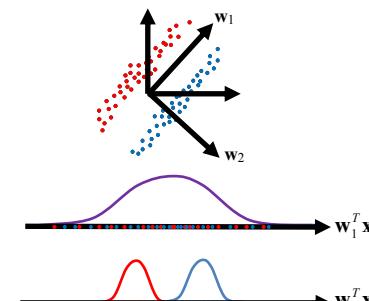


parameters in network: $10 \times 5 + 5 \times 10 = 100$
(ignoring the bias weights)

17

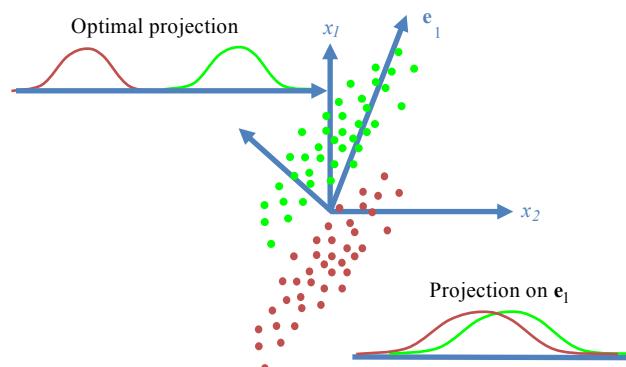
Limitations with PCA

Variance is not always the most important goal!



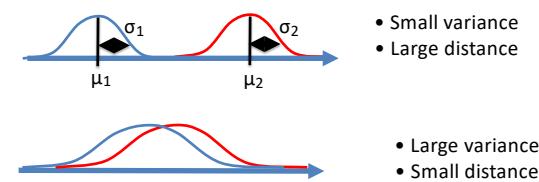
18

Optimal projection for separation of two clusters



19

Class separability



Goal: minimize variance and maximize distance.

20

Linear Discriminant Analysis (LDA)

a.k.a. Fishers Linear Discriminant (FLD)

- Minimize variance
- Maximize distance

$$\text{Maximize: } \varepsilon(\mathbf{w}) = \frac{(\mu_1 - \mu_2)^2}{\sigma_1^2 + \sigma_2^2}$$

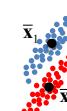
21

LDA – Cost function

$$\varepsilon = \frac{(\mu_1 - \mu_2)^2}{\sigma_1^2 + \sigma_2^2}$$

Distance:

$$\mu(\mathbf{w}) = \frac{1}{n} \sum_{i=1}^n \mathbf{w}^T \mathbf{x}_i = \mathbf{w}^T \left(\frac{1}{n} \sum_{i=1}^n \mathbf{x}_i \right) = \mathbf{w}^T \bar{\mathbf{x}}$$



$$(\mu_1(\mathbf{w}) - \mu_2(\mathbf{w}))^2 = (\mathbf{w}^T (\bar{\mathbf{x}}_1 - \bar{\mathbf{x}}_2))^2 = \mathbf{w}^T \underbrace{(\bar{\mathbf{x}}_1 - \bar{\mathbf{x}}_2)(\bar{\mathbf{x}}_1 - \bar{\mathbf{x}}_2)^T}_{\mathbf{M}} \mathbf{w} = \mathbf{w}^T \mathbf{M} \mathbf{w}$$

Variance:

$$\sigma^2(\mathbf{w}) = \frac{1}{n} \sum_{i=1}^n \left(\mathbf{w}^T (\mathbf{x}_i - \bar{\mathbf{x}}) \right)^2 = \dots = \mathbf{w}^T \underbrace{\left(\frac{1}{n} \sum_{i=1}^n (\mathbf{x}_i - \bar{\mathbf{x}})(\mathbf{x}_i - \bar{\mathbf{x}})^T \right)}_{\mathbf{C}} \mathbf{w} = \mathbf{w}^T \mathbf{C} \mathbf{w}$$

$$\sigma_1^2(\mathbf{w}) + \sigma_2^2(\mathbf{w}) = \mathbf{w}^T \mathbf{C}_1 \mathbf{w} + \mathbf{w}^T \mathbf{C}_2 \mathbf{w} = \mathbf{w}^T \underbrace{\mathbf{C}_{tot}}_{\mathbf{C}} \mathbf{w}$$

*Exercise:
Complete all the steps!*

22

LDA – Solution

$$\varepsilon(\mathbf{w}) = \frac{(\mu_1 - \mu_2)^2}{\sigma_1^2 + \sigma_2^2} = \frac{\mathbf{w}^T \mathbf{M} \mathbf{w}}{\mathbf{w}^T \mathbf{C}_{tot} \mathbf{w}}$$

Compare with PCA!

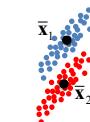
This form is called a Rayleigh quotient,
which is maximized by the largest eigenvector to the
generalized eigenvalue problem $\mathbf{C}_{tot} \mathbf{w} = \lambda \mathbf{M} \mathbf{w}$!

$$\text{Simplification: } \mathbf{M} \mathbf{w} = (\bar{\mathbf{x}}_1 - \bar{\mathbf{x}}_2) \underbrace{(\bar{\mathbf{x}}_1 - \bar{\mathbf{x}}_2)^T \mathbf{w}}_{\text{Some scalar } K} = K (\bar{\mathbf{x}}_1 - \bar{\mathbf{x}}_2)$$

Some scalar K

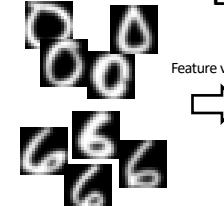
$$\mathbf{w} \sim \mathbf{C}_{tot}^{-1} (\bar{\mathbf{x}}_1 - \bar{\mathbf{x}}_2)$$

Scaling of \mathbf{w} not important!



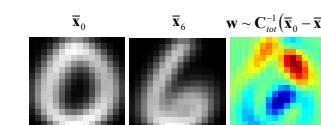
23

LDA - Example

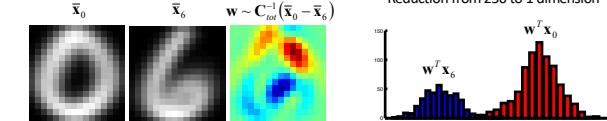


Feature vectors
→

$$\mathbf{x}_0 = \begin{pmatrix} x_1 \\ \vdots \\ x_{256} \end{pmatrix}_{0\text{-digits}} \quad \mathbf{x}_6 = \begin{pmatrix} x_1 \\ \vdots \\ x_{256} \end{pmatrix}_{6\text{-digits}}$$



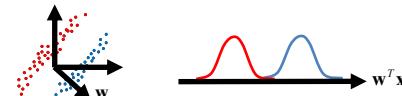
Reduction from 256 to 1 dimension



24

LDA - Summary

- Projection direction w on which two classes are maximally separated

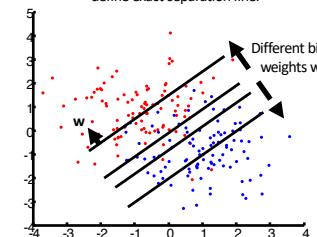


- Closed form solution, no parameters to set – easy to calculate.
- Components of w give the importance of each feature in x .
- Can be used as a classifier!

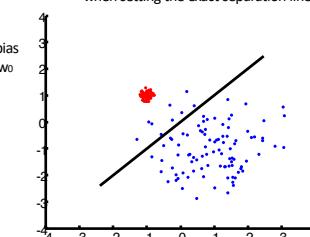
25

LDA as a classifier

Handles overlapping classes.
Must decide the bias weight to define exact separation line.



Approximates the training data with a single distribution. Loose some precision when setting the exact separation line.



26

LDA as a classifier

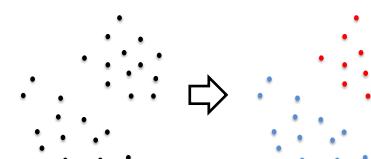
- Easy to use!
 - Closed form solution
 - Fast to calculate
 - No dependency on initialization
 - No step length to choose
 - No local optima
 - No parameters to set

Very useful as a first classifier to try and as a benchmark!

27

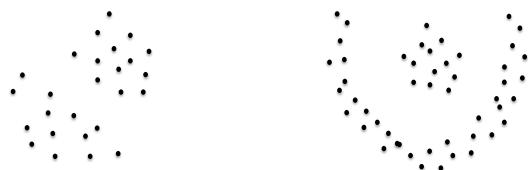
Categorization

- Categorization and grouping of objects based on similar properties is an important functionality in learning and knowledge representation.
- In the machine learning area, this is usually referred to as ***clustering***.



What describes a cluster?

- Distances to other points?
- Connectivity?
- Different definitions lead to different algorithms.



k-Means algorithm

- Assume k clusters (user input).
- Represent each cluster with a mean prototype vector \mathbf{p}_j at the cluster center.
- A data point belongs to the cluster with the closest prototype vector (Euclidian distance).



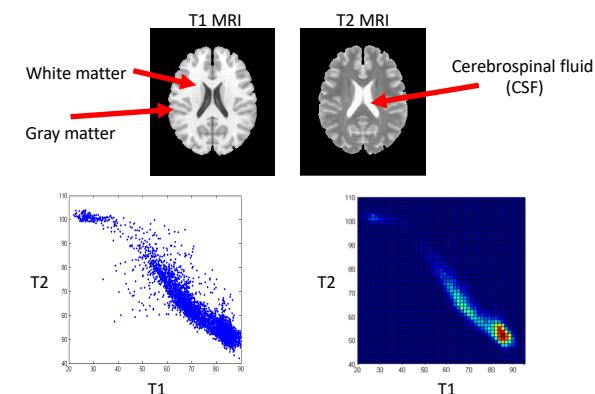
k-Means algorithm, cont.

1. Start with k random prototype vectors \mathbf{p}_j
2. Iterate:
 1. Assignment: Assign each data vector \mathbf{x}_k to the closest prototype vector \mathbf{p}_j . Denote the set of data vectors assigned to cluster \mathbf{p}_j by S_j .
 2. Update all prototype vectors to the mean of the clusters:

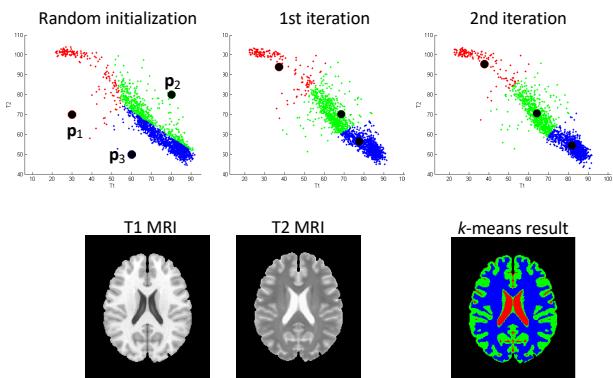
$$\mathbf{p}_j = \frac{1}{|S_j|} \sum_{k \in S_j} \mathbf{x}_k$$

Number of elements in S

k-Means - Example



k-Means - Example



k-Means - Discussion

- Must specify k
- Tries to minimize the cost function

$$\varepsilon(\mathbf{p}_1, \dots, \mathbf{p}_k, S_1, \dots, S_k) = \sum_{i=1}^k \sum_{\mathbf{x}_j \in S_i} \|\mathbf{x}_j - \mathbf{p}_i\|^2$$

- Note that this function is not differentiable as the S_i :s are discrete sets, i.e., we cannot do gradient descent.

Expectation Maximization (EM)- Intro

$$\varepsilon(\mathbf{p}_1, \dots, \mathbf{p}_k, \underbrace{S_1, \dots, S_k}_{\text{Unknown class labels}}) = \sum_{i=1}^k \sum_{\mathbf{x}_j \in S_i} \|\mathbf{x}_j - \mathbf{p}_i\|^2$$

Assume we know S_1, \dots, S_k !

$$\frac{\partial \varepsilon}{\partial \mathbf{p}_i} = \frac{\partial}{\partial \mathbf{p}_i} \left(\sum_{\mathbf{x}_j \in S_i} \|\mathbf{x}_j - \mathbf{p}_i\|^2 \right) = -2 \sum_{\mathbf{x}_j \in S_i} (\mathbf{x}_j - \mathbf{p}_i) = 2|S_i| \mathbf{p}_i - 2 \sum_{\mathbf{x}_j \in S_i} \mathbf{x}_j$$

$$\frac{\partial \varepsilon}{\partial \mathbf{p}_i} = 0 \quad \rightarrow \quad \mathbf{p}_i = \frac{1}{|S_i|} \sum_{\mathbf{x}_j \in S_i} \mathbf{x}_j \quad \text{Mean vector!}$$

Expectation Maximization (EM)- Intro

$$\varepsilon(\underbrace{\mathbf{p}_1, \dots, \mathbf{p}_k}_{\text{Continuous parameters}}, S_1, \dots, S_k) = \sum_{i=1}^k \sum_{\mathbf{x}_j \in S_i} \|\mathbf{x}_j - \mathbf{p}_i\|^2$$

Now, assume we know $\mathbf{p}_1, \dots, \mathbf{p}_k$!

Each \mathbf{x}_j independently contributes a distance $\|\mathbf{x}_j - \mathbf{p}_i\|$ to ε

Obvious that ε is minimized if each \mathbf{x}_j is assigned to the set S associated with the closest \mathbf{p} !

k-Means algorithm!!

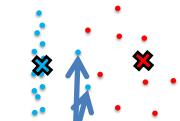
Expectation Maximization

- The k -means algorithm is a special case of a general optimization approach called *Expectation Maximization* (EM).
- EM can be used when we want to estimate model parameters (the prototypes \mathbf{p}_i), but for each data sample \mathbf{x}_j , there is a hidden/missing parameter (the class labels S_j).
- EM iterates between optimizing the hidden parameters and the model parameters.

Mixture of Gaussians (MoG) clustering

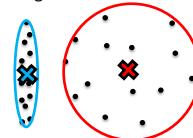
Another application of EM!

k-means



May not be correctly clustered!

Idea: Represent each cluster using a Gaussian distribution



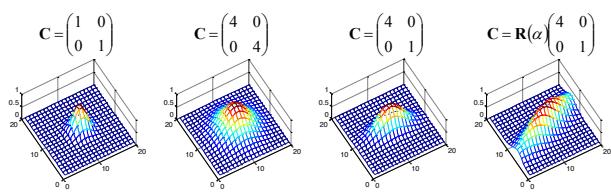
Problem: Each data sample \mathbf{x}_j belongs to one of k Gaussian distributions $N(\mathbf{p}_i, \mathbf{C}_i)$, $i=1..k$.

Find the sets S_i of samples that belong to distribution $N(\mathbf{p}_i, \mathbf{C}_i)$ and the mean \mathbf{p}_i and covariance matrix \mathbf{C}_i .

The Gaussian distribution

$$f(\mathbf{x}; \mathbf{p}_i, \mathbf{C}_i) = \frac{1}{(2\pi)^{d/2} |\mathbf{C}_i|^{1/2}} e^{-\frac{1}{2} (\mathbf{x} - \mathbf{p}_i)^T \mathbf{C}_i^{-1} (\mathbf{x} - \mathbf{p}_i)}$$

↑ Dimension ↑ Determinant of covariance matrix



Let's use EM!

Assume we know the hidden parameters S_1, \dots, S_k !

That is, we know the samples for each set, e.g., $S_1 = \{\mathbf{x}_1, \mathbf{x}_7, \mathbf{x}_{12}, \mathbf{x}_{13}\}$

We can then estimate the mean \mathbf{p}_i and covariance \mathbf{C}_i using standard estimation for the Gaussian:

$$\mathbf{p}_i = \frac{1}{|S_i|} \sum_{k \in S_i} \mathbf{x}_k$$

$$\mathbf{C}_i = \frac{1}{|S_i|} \sum_{k \in S_i} (\mathbf{x}_k - \mathbf{p}_i)(\mathbf{x}_k - \mathbf{p}_i)^T$$

Let's use EM, cont!

Assume now that we know the Gaussian distribution parameters, i.e., we know all distributions

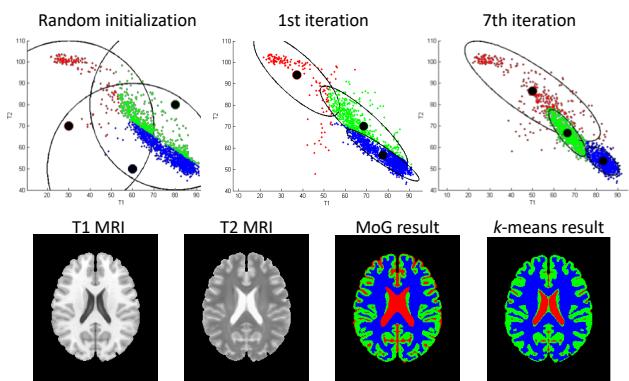
$$f(\mathbf{x}; \mathbf{p}_i, \mathbf{C}_i) = \frac{1}{(2\pi)^{d/2} |\mathbf{C}_i|^{1/2}} e^{-\frac{1}{2} (\mathbf{x} - \mathbf{p}_i)^T \mathbf{C}_i^{-1} (\mathbf{x} - \mathbf{p}_i)}$$

Let the hidden parameter S_i be the set of all data samples for which $f(\mathbf{x}; \mathbf{p}_i, \mathbf{C}_i)$ is larger than for all other distributions. That is, each data sample is assigned to the Gaussian distribution **to which it most likely belongs!**

Mixture of Gaussians - Algorithm

1. Start with k random Gaussians $(\mathbf{p}_j, \mathbf{C}_j)$
 2. Iterate:
 1. Assign each sample \mathbf{x}_i to the most likely Gaussian
- $$\max_j f(\mathbf{x}_i; \mathbf{p}_j, \mathbf{C}_j) = \frac{1}{(2\pi)^{d/2} |\mathbf{C}_j|^{1/2}} e^{-\frac{1}{2} (\mathbf{x}_i - \mathbf{p}_j)^T \mathbf{C}_j^{-1} (\mathbf{x}_i - \mathbf{p}_j)}$$
- Denote the set of samples assigned to Gaussian j by S_j .
2. Update all Gaussian means and covariances:
- $$\mathbf{p}_j = \frac{1}{|S_j|} \sum_{k \in S_j} \mathbf{x}_k$$
- $$\mathbf{C}_j = \frac{1}{|S_j|} \sum_{k \in S_j} (\mathbf{x}_k - \mathbf{p}_j)(\mathbf{x}_k - \mathbf{p}_j)^T$$

Mixture of Gaussians - Example



Summary of k -Means and MoG clustering

- Must choose the number of clusters k manually.
- Different initializations may give different results.
- May converge to degenerate solutions, e.g., empty clusters.
- MoG allows for elliptic cluster shapes with different sizes.